

B V RAVIKIRAN

ravikiranbv1806@gmail.com +91-9025465892 linkedin.com/in/ravi-kiran-43528b20b

Professional Summary

Backend Engineer (AI/LLM Systems) with 2+ years building production Python services and Agentic AI workflows for enterprise-scale platforms. Specialized in RAG architectures, LLM orchestration (GPT, Anthropic), and FastAPI-based backend systems currently automating multi-step report migration workflows at Ascendion, reducing manual effort through AI-driven conversion. Skilled in end-to-end delivery: API design, containerization (Docker), cloud deployment (Azure, AWS), and CI/CD. Seeking to bring AI-integrated backend engineering expertise to [target role/domain]

Professional Experience

Associate Engineer, Ascendion
Insights Modernization Studio (IMS)

07/2024 – Present
Chennai, India

- Developed scalable backend services using Python, FastAPI, and PostgreSQL for an enterprise report migration and modernization platform, combining manual Python-based integrations with Agentic AI-driven automation.
- Designed and built 45+ AI agents and 15 Agentic AI workflows using GPT, Claude, and OpenAI APIs, reducing manual coding effort by 50% for multi-step migration tasks.
- Built and integrated REST APIs, including Microsoft Fabric REST APIs, to automate metadata extraction and support end-to-end report migration workflows.
- Built automated Report migration workflows using LLMs..
- Integrated Agentic AI and LLM workflows into the FastAPI platform to automate report transformation, cutting manual migration effort and accelerating delivery timelines.
- Containerized the application using Docker and deployed to Azure App Service, supporting CI/CD-driven collaborative development via Git and Azure Repos.

Education

B.E. in Computer Science, Sri Krishna College of Engineering and Technology 2020 – 2024
CGPA: 8.08

Technical Skills

Backend Development
Python | FastAPI | Docker | PostgreSQL | Claude

Cloud
AWS | Azure

AI Engineering & Developer Tools
LLM Orchestration | Prompt Engineering | OpenAI | RAG | GitHub Copilot | Agentic AI | Langchain

Databases
PostgreSQL | SQL

Projects

DocuMind AI , Smart Document Copilot & Retrieval Assistant
Built a full-stack RAG-based PDF chatbot with FastAPI, React, and PostgreSQL (pgvector), tested with 50+ PDF documents for semantic search and Q&A, using Groq-hosted Llama 3.1 8B for low-latency inference and Sentence Transformers for embedding generation via LangChain's text splitting utilities. Architected as 3 containerized microservices (backend, frontend, database) via Docker Compose, deployed to EC2-provisioned compute nodes for orchestration and scaling. Reduced Docker image size by optimizing for CPU-only inference.

ToDoList , Fast API focused backend App
Developed a full stack Todo application using FastAPI and SQLAlchemy, building 12+ REST API endpoints with secure cookie-based session authentication and PBKDF2-SHA256 password hashing via Passlib. Implemented user registration, login, and per-user authorization checks ensuring data isolation across all task CRUD operations. Added structured logging for request tracing, authentication events, and error handling to support debugging and monitoring. Containerized the application with Docker and deployed to AWS EC2, with database-agnostic configuration supporting both SQLite (local) and PostgreSQL (server) via environment-driven connection strings.

Stock Assistant , RAG pipeline for stock research, powered by LangChain, microservices, and real-time price + news retrieval
A microservices-based RAG stock research assistant with three independently deployable FastAPI/script services (ingestion, retrieval, chat) sharing a Postgres+pgvector database. It combines structured price data (yfinance) and unstructured news (NewsAPI, embedded via sentence-transformers) using a LangChain LCEL router that decides which source(s) a question needs, then generates grounded answers via Groq. Chat-service adds JWT-authenticated users with persistent, per-conversation memory stored in Postgres, and multi-turn follow-up resolution via query condensing. Built for containerized local development (Docker Compose) and cloud deployment (AWS ECS/EKS), with structured logging designed for CloudWatch aggregation.

Certifications

- Microsoft Azure Fundamentals (AZ-900)
- Claude AI Frequency and Framework
- Claude Intoduction to MCP